

THU SUMMER SCHOOL 2025

CHANTAL DEUSCH | JEREMY DIEM | JAN GASCHLER | SERHAT GÜREL | PAULINA PYCZOT | VALENTIN TALMON-L'ARMÉE

Image Detection with YOLO AI on a Raspberry Pi



Agenda

1. Setup Raspberry Pi
2. Licence Plate Recognition
3. Face Recognition



Visual Process Structure



Raspberry Pi

Taking Pictures



Detect Plate



Reading Plate



Live Stream



Detect Person
and Face




Learn Face



Set Up Raspberry Pi

To set up the Raspberry Pi, you will need the following physical components:

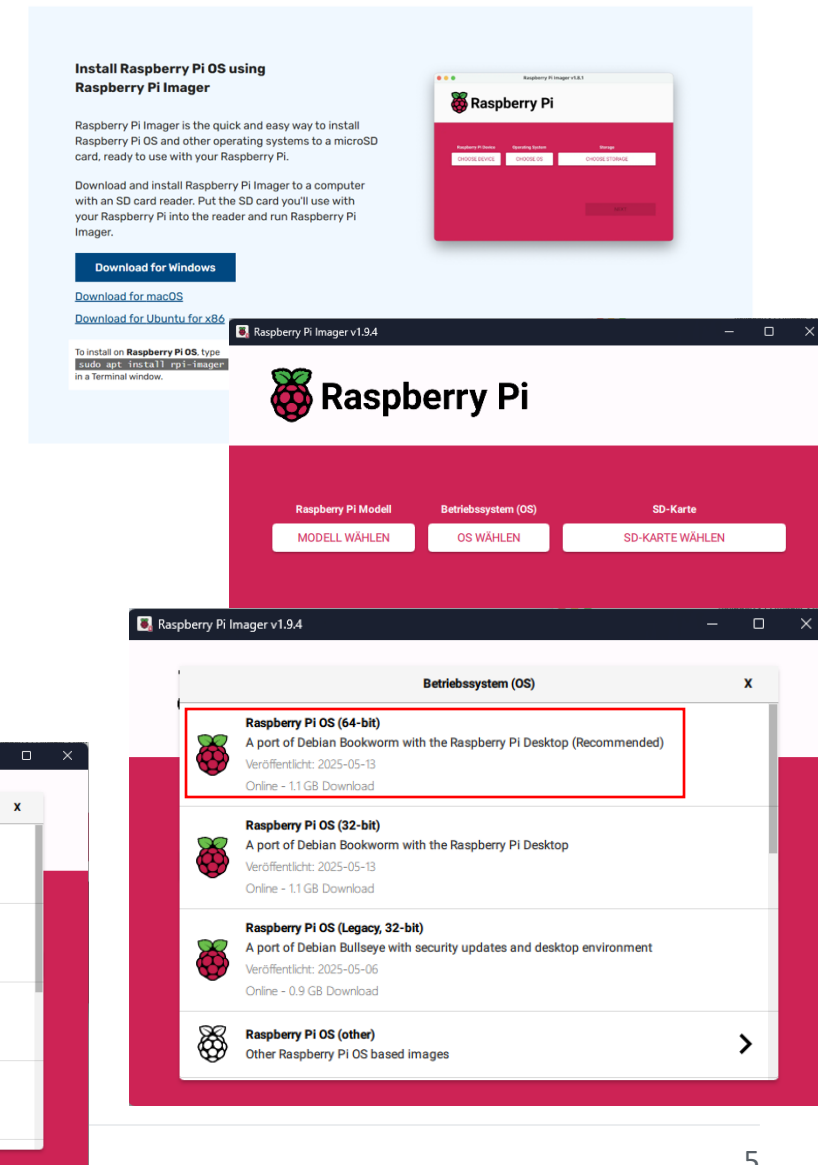
- SDHC card
- USB-compatible camera
- SDHC Adapter
- External PC/Laptop
- Ethernet Switch/Internet connection



Preparing the SDHC Card

Formatting the SDHC Card and installing Raspberry Pi OS

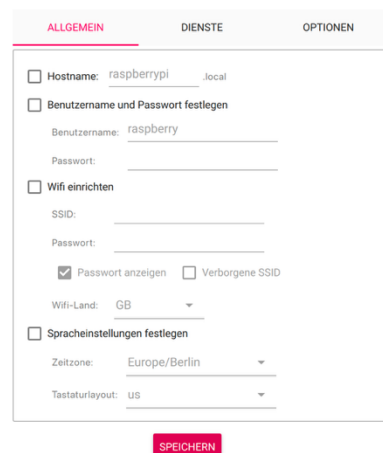
1. First, you need to download and install the Raspberry Pi Imager Software on your Computer:
<https://www.raspberrypi.com/software/>
2. Once installed, you'll be greeted with the options to choose a model, an OS and an SDHC Card
3. For the model, we'll choose **Raspberry Pi 4** and for the Operating System (OS) **Raspberry Pi OS (64-bit)**



Preparing the SDHC Card

Formatting the SDHC Card and installing Raspberry Pi OS

1. Then you need to select the inserted SDHC Card where the OS should be installed by clicking “Choose SD Card”
2. Finally, you need to click on continue and then “Edit settings” and then “Services” and then tick the “Enable SSH” option to enable the ssh server for the connection later
3. After this, hit save, yes and yes



ALLGEMEIN DIENSTE OPTIONEN

☐ Hostname: raspberrypi .local

☐ Benutzername und Passwort festlegen

Benutzername: raspberry

Passwort: _____

☐ Wifi einrichten

SSID: _____

Passwort: _____

☒ Passwort anzeigen ☐ Verborgene SSID

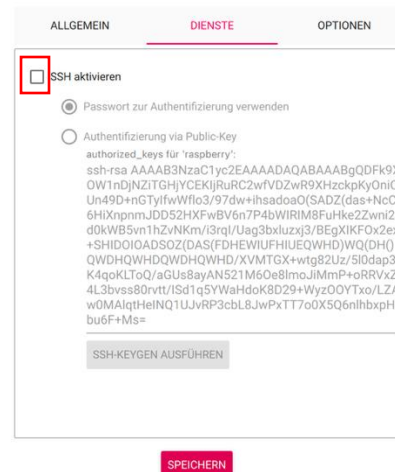
Wifi-Land: GB

☐ Spracheinstellungen festlegen

Zeitzone: Europe/Berlin

Tastaturlayout: us

SPEICHERN



ALLGEMEIN DIENSTE OPTIONEN

☒ SSH aktivieren

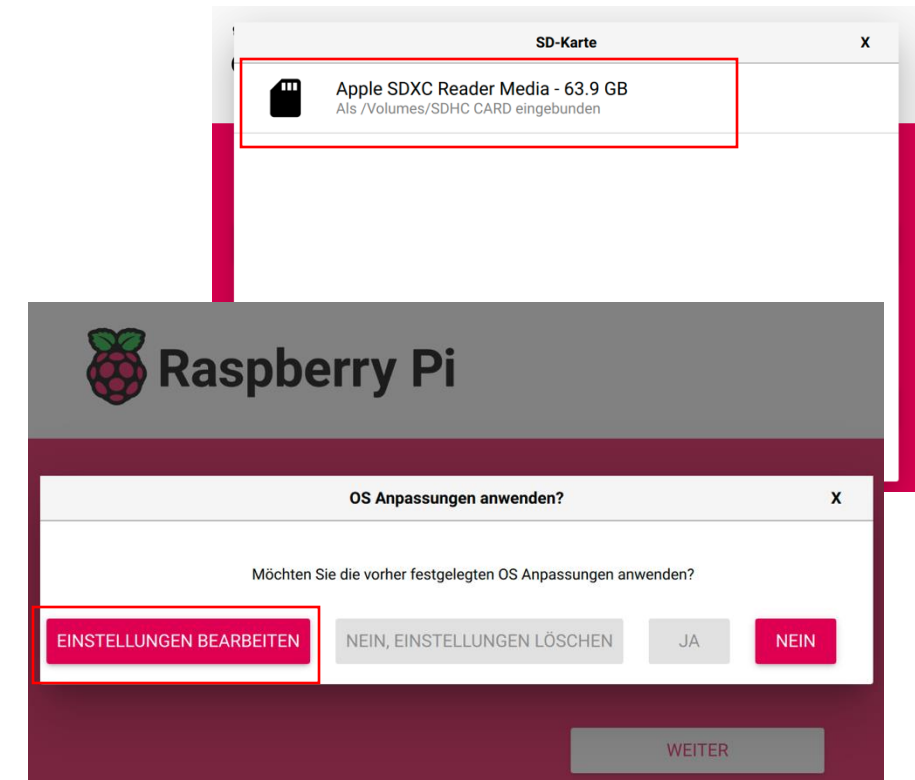
☒ Passwort zur Authentifizierung verwenden

☐ Authentifizierung via Public-Key

authorized_keys für 'raspberrypi':

```
ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQGDQFk9XOW1nDjNZITGhYCEKjRuRC2wFVDZwR9XHczkPyOniCUn49D+nGTylfwWflo3/97dw+hsadoaO(SADZ(das+NcC6HiXnpnmJDD52HXFwBV6n7P4bWIRIM8FuHkeZ2wni2d0kWB5m1hZvNkm/i3rq/Uag3bXluzxj3/BEgKIKFOx2ex+SHdOioQADSQZ(DAS(FDHEWUHFHUEQWHD)WQ(DH)QWDHQWHDQWHDQWHD/XVMTGX+wrtg8Zuz/5l0dap3K4qoKLT0/aGUs8ayAN521M60e8lmeJlMmP+oRRVxZ4L3bvss80rvtt/ISd1q5YWaHdoK8D29+WyzO0YTxo/LZA w0MAIqtHelNQ1UJvRP3cbl8JwPxTT7o0X5Q6nlhxpHbu6F+Ms=
```

SPEICHERN



SD-Karte X

Apple SDXC Reader Media - 63.9 GB
Als /Volumes/SDHC CARD eingebunden

Raspberry Pi

OS Anpassungen anwenden? X

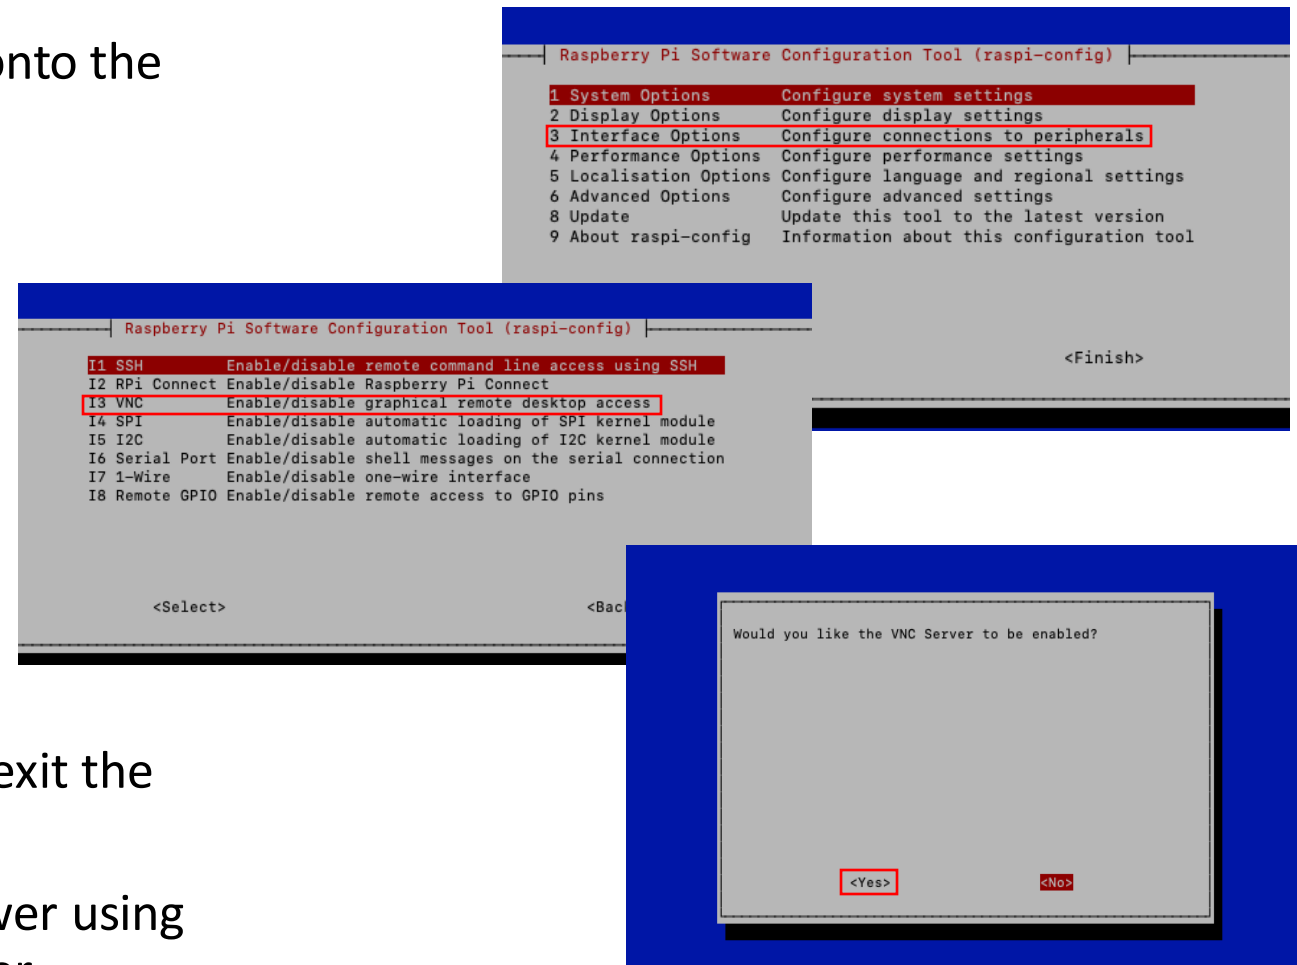
Möchten Sie die vorher festgelegten OS Anpassungen anwenden?

EINSTELLUNGEN BEARBEITEN NEIN, EINSTELLUNGEN LÖSCHEN JA NEIN

WEITER

Enable vnc using ssh

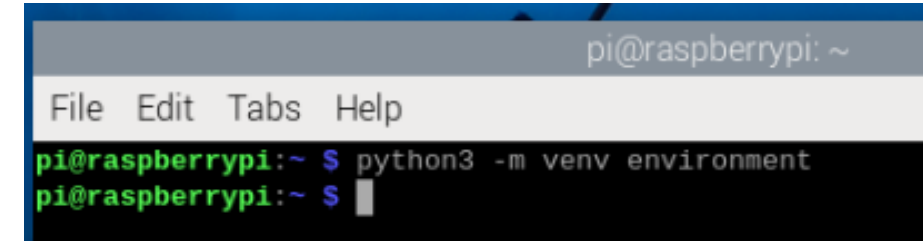
1. Use the IP Address of the Raspberry Pi to ssh onto the Pi using **ssh pi@[IP]**
Password: *raspberry*
2. Then use **sudo raspi-config** to enter the Raspberry Pi Configuration menu
3. In the menu navigate to *Interface Options* and *enter*
4. In the *Interface options* navigate to *VNC* and *enter*
5. In the popup select *Yes* and *enter*
6. Now the VNC server is ready to use, you can exit the configuration menu now
7. You can connect using your favorite VNC Viewer using **[IP of the Pi]:5900** and the already known user credentials



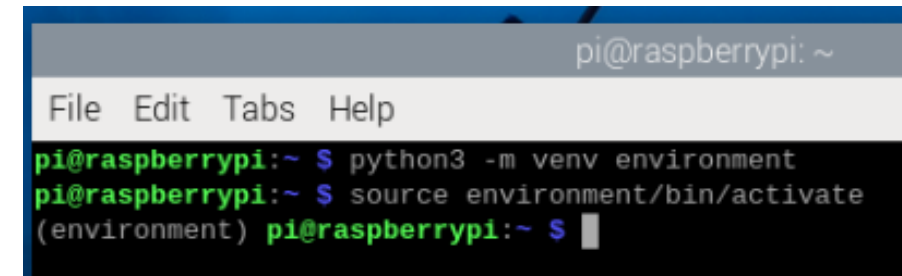
Preparing the venv environment

Creating and using a venv environment

1. Using the Raspberry Pi Terminal you can create and use a Python Virtual Environment
2. First, type **python3 -m venv environment** in the terminal
3. Then enable the created venv environment using **source environment/bin/activate**
4. Now we can go over to prepare the libraries using our Script



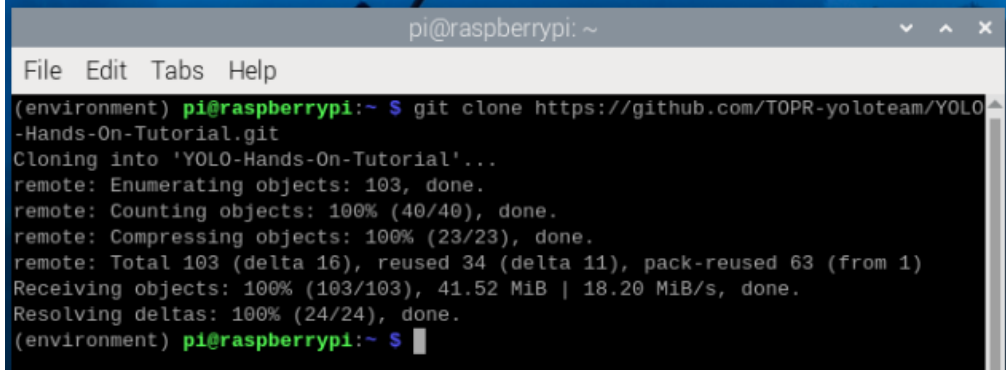
```
pi@raspberrypi: ~  
File Edit Tabs Help  
pi@raspberrypi:~ $ python3 -m venv environment  
pi@raspberrypi:~ $
```



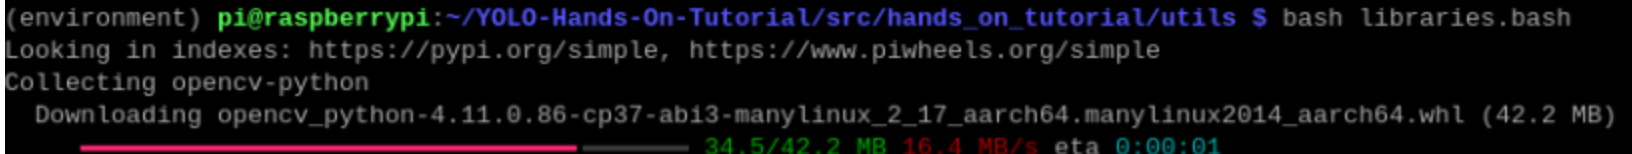
```
pi@raspberrypi: ~  
File Edit Tabs Help  
pi@raspberrypi:~ $ python3 -m venv environment  
pi@raspberrypi:~ $ source environment/bin/activate  
(environment) pi@raspberrypi:~ $
```


Preparing the needed libraries

1. First, you need to clone the GitHub Repository using:
git clone https://github.com/TOPR-yoloteam/YOLO-Hands-On-Tutorial.git
2. Then you need to go into the utils folder using:
cd YOLO-Hands-On-Tutorial/src/hands_on_tutorial/utils
3. Now execute the bash file using:
bash libraries.bash
4. Now all the needed libraries and everything else will be automatically installed and ready to use!



```
pi@raspberrypi: ~  
File Edit Tabs Help  
(environment) pi@raspberrypi:~ $ git clone https://github.com/TOPR-yoloteam/YOLO-Hands-On-Tutorial.git  
Cloning into 'YOLO-Hands-On-Tutorial'...  
remote: Enumerating objects: 103, done.  
remote: Counting objects: 100% (40/40), done.  
remote: Compressing objects: 100% (23/23), done.  
remote: Total 103 (delta 16), reused 34 (delta 11), pack-reused 63 (from 1)  
Receiving objects: 100% (103/103), 41.52 MiB | 18.20 MiB/s, done.  
Resolving deltas: 100% (24/24), done.  
(environment) pi@raspberrypi:~ $
```



```
(environment) pi@raspberrypi:~/YOLO-Hands-On-Tutorial/src/hands_on_tutorial/utils $ bash libraries.bash  
Looking in indexes: https://pypi.org/simple, https://www.piwheels.org/simple  
Collecting opencv-python  
  Downloading opencv_python-4.11.0.86-cp37-abi3-manylinux_2_17_aarch64.manylinux2014_aarch64.whl (42.2 MB)  
 34.5/42.2 MB 16.4 MB/s eta 0:00:01
```

License Plate Recognition



License Plate Recognition

Helpful tips

- › To see changes better, use the following code after changing an image
- › Code to be written is marked with #TODO X in the source code
- › Google loves you, ChatGPT doesn't <3

```
cv2.imshow( winname: "image", image)  
cv2.waitKey(0)
```

```
#TODO 1
```

Taking Pictures

- › To recognize license plates, we need images. These images must be captured first.
We do this using the *Image* class in the *make_images.py* file.
 - Run the script to start capturing images from the webcam



License Plate Recognition

Detect License Plate

- › In order to be able to read the text from license plates later, we must first recognize the license plate as such
- › **Task 1: Modify the Python code to detect a license plate and draw its corresponding bounding box**
- › *Hint:* The Ultralytics docs homepage provides information on how to use a YOLO model :')



License Plate Recognition

Detect License Plate

- › To achieve the best OCR results, we now want to save the license plate as a single image
- › **Task 2: Save the license plate as a single image in a separate folder named "license_plates"**
- › *Hint:* Use the plates bounding box coordinates!



License Plate Recognition

Read License Plate

- › Colored images can cause problems in character recognition, therefore we need to preprocess the input image
- › **Task 1:** Preprocess Image on color, threshold, dilation and contours
- › *Hint:* Use the OpenCV library for image preprocessing. For contours use the given method "find_and_sort_contours"



Read License Plate

› If all steps are done correctly, start with the OCR!

› **Task 2:** Extract text from the contours and print the results



› *Hint:* What does "extract_text_from_contours" return?



```
Image: image_0_0.png
```

```
Text: ULDT805 | Probabilities: 86.33
```

Face Recognition



Please choose between:


- YOLO
 - › In this tutorial, we explore how the YOLOv8n-face model can be used to detect human faces in real time
 - › The YOLOv8n-face architecture is fine-tuned specifically for faces, combining high detection accuracy with a relatively small model size (only 1 object class)
 - › Although **YOLO** is highly optimized for GPU acceleration and real-time performance, it tends to lag more than **MediaPipe**, especially on resource-constrained devices like a Raspberry Pi
- MediaPipe
 - › A framework by Google providing lightweight models, particularly FaceMesh for detailed facial landmark detection.
 - › Often highly performant, even on resource-constrained devices like the Raspberry Pi.
 - › Focuses on extracting facial features (landmarks) which can then be used for recognition by comparing their relative positions.

If you want to solve the MediaPipe exercises, please skip to page **33**

YOLO

Task 1:

- › Implement **face detection** based on YOLO.
- › Detect **faces** and draw bounding boxes around them.



YOLO Task 1 tips

Task 1.1:

- › Ensure the `self.model_path` correctly points to the "yolov8n-face.pt" file within the project structure.

Task 1.2:

- › You run detection by calling the model object directly: `results = self.model(...)`.
- › Pass the video frame as the main argument to the model.
- › **Crucial:** Use the `classes=[0]` argument to tell the model to *only* detect faces (which is class 0 in this specific face model).

Task 1.3:

- › The results object contains detections. You often need to loop through results, then `result.bboxes`, then each box.
- › The coordinates are typically stored within a box object, often under the attribute `.xyxy[0]`.
- › The coordinates might be floating-point numbers. Use `map(int, ...)` to convert them to integers before drawing.
 - › *Example:* `x1, y1, x2, y2 = map(int, box.xyxy[0])`

YOLO Task 1 tips

Task 1.4:


- › You can use OpenCV's drawing function: `cv2.rectangle(...)`.
- › Remember the required arguments: `cv2.rectangle(image, pt1, pt2, color, thickness)`.
 - › *pt1* is the top-left corner (x1, y1).
 - › *pt2* is the bottom-right corner (x2, y2).
 - › *Color* is a BGR tuple, e.g., (0, 255, 0) for green.
 - › *Thickness* is an integer, e.g., 2.

If you want to solve the MediaPipe exercises, please skip to page **33**

YOLO

Task 2:

- › Extend the system by adding face recognition.
Person detection is **already provided**.
Implement recognizing known faces and mark them in the image.



YOLO Task 2 tips

Task 2.1:

- › Make sure you're reading image files from the correct directory (`self.faces_dir`).
- › Use `face_recognition.load_image_file()` to load each image.
- › Get the face encoding using `face_recognition.face_encodings(loaded_image)[0]`.
 - › Note the `[0]` assumes one face per known image file.
- › Store the encoding in `self.known_face_encodings` and the name (from filename) in `self.known_face_names`. Keep these lists synchronized!

Task 2.2:

- › Get the `face_location` passed to the `save_face` function.
- › Crop the face from the stored `self.current_frame` using these coordinates: `face_image = self.current_frame[top:bottom, left:right]`.
- › Consider expanding the coordinates slightly before cropping to get more context, but be careful not to go outside the frame boundaries (`max(0, ...)`, `min(width, ...)`).
- › Save the `face_image` using `cv2.imwrite()` into `self.faces_dir`. Give it a filename based on the entered name.
- › **Very Important:** After saving, immediately call `self.load_known_faces()` again so the system recognizes the newly added person right away.

YOLO Task 2 tips

Task 2.3, 2.5., 2.6, 2.7:

- › Use the self.state variable ("normal" or "entering_name") to control behavior.
- › In mouse_callback: If state is "normal", check if the click (x, y) is inside any stored button area. If yes, change self.state to "entering_name" and store which face was clicked (self.selected_face_loc).
- › In the main loop (run): If self.state is "entering_name", handle keyboard input (cv2.waitKey(1)):key == 13 (Enter): Save the face if self.current_text is not empty, then switch back to "normal" state.
 - › key == 27 (Esc): Switch back to "normal" state.
 - › key == 8 (Backspace): Remove last character from self.current_text.
 - › Printable characters: Append chr(key) to self.current_text.
- › Draw the text input UI (draw_text_input) and highlight the selected face only when in the "entering_name" state.

YOLO Task 2 tips

Task 2.4:

- › You need face locations and encodings for the *current* video frame first.
- › Use `face_recognition.face_locations()` and `face_recognition.face_encodings()`.
- › **Remember:** Convert the OpenCV frame (BGR) to RGB before passing it to face_recognition functions: `rgb_frame = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)`.
- › Use `face_recognition.compare_faces(self.known_face_encodings, encoding_to_check)` to get potential matches.
- › Use `face_recognition.face_distance(self.known_face_encodings, encoding_to_check)` to find the distance for each comparison. Lower distance means a closer match.
- › Find the best match using `np.argmin(face_distances)` on the distances array.
- › Verify the best match: Check if `matches[best_match_index]` is True.

YOLO Task 2 tips

Task 2.4:


- › Draw the name label using `cv2.putText()`.
- › Draw the "Learn Face" button (using `cv2.rectangle` and `cv2.putText`) **only** if the face is "Unknown" and the system is in the "normal" state.
- › For the button click to work later, store the button's screen coordinates *and* the corresponding face location (top, right, bottom, left) together (e.g., in `self.button_area`).

If you want to solve the YOLO exercises, please go back to page 25

MediaPipe

Task 1:

- › Implement **face detection** based on MediaPipe.
- › Face recognition is not required in this task.
- › Detect **faces** and draw bounding boxes around them.



MediaPipe Task 1 tips

Task 1.1:

- › Create the FaceMesh instance using `mp.solutions.face_mesh.FaceMesh(...)`.
- › You can set parameters like `max_num_faces=...` and `min_detection_confidence=...` during initialization (check MediaPipe documentation for details).

Task 1.2:

- › **Important:** MediaPipe expects images in RGB format, but OpenCV captures in BGR. Convert the frame: `image_rgb = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)`.
- › Call the `process()` method on your `self.face_mesh` object, passing the `image_rgb`.

Task 1.3:


- › Use OpenCV's `cv2.rectangle(image, pt1, pt2, color, thickness)`. `pt1` should be the top-left corner: (left, top). `pt2` should be the bottom-right corner: (right, bottom).
- › Choose a color (BGR tuple like (0, 255, 0) for green) and thickness.

If you want to solve the YOLO exercises, please go back to page 25

MediaPipe

Task 2:

- › Extend the system by adding face recognition.
- › Person detection is **already provided**.
- › Implement recognizing known faces and mark them in the image.



MediaPipe Task 2 tips

Task 2.1:

- › You don't need all 400+ landmarks from FaceMesh. Use the provided `key_landmarks_indices` list to select a smaller, specific set (e.g., 50 landmarks).
- › Loop through `key_landmarks_indices`. For each index `idx`, get the landmark data: `landmark = face_landmarks.landmark[idx]`.
- › Store the **normalized** coordinates `landmark.x` and `landmark.y` (values between 0.0 and 1.0).
- › Create a flat NumPy array from these coordinates (e.g., `[x1, y1, x2, y2, ..., x50, y50]`). The size must be consistent!

Task 2.2:

- › This happens in the `compare_landmarks` function.
- › You need to compare the detected landmarks array to *every* stored landmark sample for *every* known person (for `known_landmarks` in `person_landmarks_list`).
- › A good way to measure similarity is the Euclidean distance: `distance = np.linalg.norm(landmarks - known_landmarks)`.
- › Keep track of the `min_distance` found and the `best_match_index` (which person it corresponds to).

MediaPipe Task 2 tips

Task 2.3:

- › After checking all known landmarks, take the overall min_distance.
- › Compare it to self.recognition_threshold.
- › If min_distance < self.recognition_threshold, it's a match! Set is_known_face = True and get the name using best_match_index.
- › If the distance is larger, it's "Unknown" (is_known_face = False).
- › Consider calculating a confidence score based on how far below the threshold the distance is.